

INDEPENDENT ANALYSIS

The Claude Code Tools Guide

An unbiased, evidence-based evaluation of the top 9 Claude Code plugins, frameworks, and orchestration tools

March 2026 Edition

Superpowers 71,341 ★

Everything CC 50,000+ ★

UI/UX Pro Max 37,340 ★

GSD 24,700 ★

gstack **NEW** 22,134 ★

Compound Eng. 9,893 ★

Claude Skills 4,919 ★

Ralphy 2,517 ★

Plugins+Skills 1,513 ★

00

Contents

01

Executive Summary

.....

02

Tool Profiles & Classification

.....

03

Classification Matrix

.....

04

Deep Comparison

.....

05

Honest Reality Check

.....

06

Recommendations by Scenario

.....

07

Combining Tools

.....

08

Final Ranking

.....

01

Executive Summary

Nine tools were evaluated through direct repository inspection: README analysis, file structure exploration, skill/workflow content review, release history tracking, and community metrics assessment. Each repository was classified by its actual function rather than its marketing.

Use Case Recommendations

USE CASE	RECOMMENDATION
Best overall choice	Superpowers
Best for beginners	Claude Skills
Best for power users	Superpowers
Best for long / complex projects	Get Shit Done (GSD)
Best for structured / spec-driven workflows	Get Shit Done (GSD)
Best for autonomous loop workflows	Ralphy
Best for domain expertise injection	Claude Skills
Best for UI/UX design generation	UI/UX Pro Max Skill
Best for comprehensive all-in-one setup	Everything Claude Code
Best for plugin catalog & learning	Plugins Plus Skills
Best for role-based AI team workflows	gstack

Ecosystem Update: GSD has expanded with GSD-2, adding a 27-document knowledge base and advanced orchestration patterns. gstack (by Y Combinator CEO Garry Tan) introduces role-based cognitive modes and production-proven review checklists.

Author's Pick

Superpowers would be my first install. It changes how the agent thinks, not just what it does. The anti-rationalization architecture and TDD enforcement produce measurably better code. Then add **Claude Skills** as a complementary expertise layer. For long projects, add **GSD**. For UI-heavy work, add **UI/UX Pro Max**. For role-based workflows with QA and design audit, add **gstack**. The key insight: these tools occupy different layers and the best setup combines 2-3, not all 9.

Tool Profiles

Superpowers

METHODOLOGY FRAMEWORK

by Jesse Vincent (obra) • 71.3K ★ • 5,489 forks • 13 skills • 3 commands • v4.3.1 • 4 platforms

"An opinionated software development methodology delivered as composable skills. Enforces brainstorming before implementation, TDD (no production code without a failing test), systematic debugging, and rigorous code review. The distinguishing innovation is its anti-rationalization architecture."

- + Battle-tested methodology (17 months, 71K community)
- + Zero config—session-start hook auto-loads methodology
- + Cross-platform: Claude Code, Cursor, Codex, OpenCode
- + Exceptional documentation depth per skill
- Prescriptive workflow (can't skip phases easily)
- No orchestration loop or parallel execution
- Windows support is fragile (6+ documented fixes)

Get Shit Done (GSD)

ORCHESTRATION FRAMEWORK

by TACHES (gsd-build) • 24.7K ★ • 2,110 forks • 12 agents • 32 commands • v1.22.4 • 35 workflows

"A multi-agent orchestration framework purpose-built to fight context rot. Uses a lean orchestrator (~10-15% context) that dispatches subagents with fresh 200K-token windows. Features wave-based parallel execution, goal-backward verification, Nyquist test validation, and atomic git commits. Now accompanied by GSD-2."

- + Context engineering solves a real problem (context rot)
- + Full project lifecycle coverage (32 commands)
- + Wave-based parallel execution with dependency awareness
- + Goal-backward verification (checks results, not claims)
- Steep learning curve (32 commands to discover)
- Rapid churn risk (22 releases in 2 weeks, 258 open issues)
- Overkill for simple tasks

gstack

NEW

ROLE-BASED SKILL SYSTEM

by Garry Tan / Y Combinator • 22.1K ★ • 13 skills • headless browser CLI • v1.1.0 • Bun runtime

"A software factory that turns Claude Code into a virtual engineering team — CEO, Eng Manager, Designer, Staff Engineer, QA Lead, Release Engineer, and Technical Writer. Each role is a complete workflow skill with opinionated checklists, review gates, and persistent browser-based QA testing."

- + Role-based cognitive modes (CEO/Eng/Design/QA)
- + Persistent headless browser for real QA testing (~100ms)
- + Fix-First review methodology (AUTO-FIX vs ASK)
- + AI Slop detection patterns for frontend code
- + "Boil the Lake" completeness principle
- + Production-battle-tested review checklists
- Rails/Ruby-centric examples and patterns
- No multi-agent orchestration (external Conductor)
- Single-session monolithic skills (no delegation)
- Requires Bun runtime for browser CLI

Compound Engineering

PLUGIN + WORKFLOW + CLI

by EveryInc (Kieran Klaassen) • 9.9K ★ • 783 forks • 29 agents • 22 commands • v2.34 • 20 skills

"A dual-purpose project: the Compound Engineering plugin (plan → work → review → compound methodology) and a cross-platform CLI converter that translates Claude Code plugins into formats for 10 other AI tools."

- + Cross-platform converter (10 AI tools)—unique
- + Institutional knowledge capture (compound-docs)
- + Deepest feature set (29 agents, 22 commands, 20 skills)
- Rails/Ruby bias in review agents
- Large context footprint (48KB swarm skill alone)
- Single-maintainer dominated (71% of commits)

Claude Skills

SKILLS PACK / PLUGIN

by Jeff Smolinski (jeffallan) • 4.9K ★ • 308 forks • 66 skills • 365 references • v0.4.9 • 12 domains

"The broadest knowledge-injection system: 66 specialized skills across 12 domains. Uses progressive disclosure—compact SKILL.md files with deep reference files loaded on demand for token efficiency."

- + Widest domain coverage (66 skills, 12 domains)
- + Token-efficient progressive disclosure architecture
- + Multi-agent support (11 AI coding tools)
- Pre-1.0 maturity (v0.4.9)
- No orchestration or execution engine
- Primarily solo contributor (89% of commits)

Ralphy

AUTONOMOUS TASK RUNNER

by Michael Shimeles • 2.5K ★ • 313 forks • 8 engines • 5 task formats • v4.7.2 • 10+ contributors

"A standalone CLI that orchestrates AI coding agents in an autonomous loop until a PRD is fully implemented. Supports 8 AI engine backends, parallel execution with git worktrees or sandboxes."

- + Simplest mental model: ralphy --prd PRD.md
- + Engine-agnostic (8 AI backends)
- + Parallel execution with sandbox isolation
- Runs with --dangerously-skip-permissions
- Zero quality guardrails (no TDD, no review)
- Bad PRDs produce bad code at machine speed

Everything Claude Code

COMPREHENSIVE PLUGIN SYSTEM

by Affaan Mustafa (affaan-m) • 50K+ ★ • 6K+ forks • 58 skills • 13 agents • v1.7.0 • 35 commands

"A comprehensive, battle-tested productivity system built on 10+ months of intensive daily production use. Includes 58 skills, 13 specialized agents, 35 commands, a mature hook system, multi-language rules architecture, and security-first design with AgentShield integration."

- + 992 internal validation tests—most thoroughly tested
- + Security-first: AgentShield, 1,282 security tests
- + Multi-language rules (common/ + language-specific/)
- "Kitchen sink" risk—58 skills + 35 commands
- Hook debugging is opaque (10s timeouts)
- Node.js dependency for hooks and scripts

Plugins Plus Skills

MARKETPLACE + LEARNING HUB

by Jeremy Longshore • 1.5K ★ • 182 forks • 315 plugins • 1,537 skills • v4.15 • 15 tutorials

"The largest single-repo Claude Code marketplace: 315 plugins across 20+ categories with 1,537 validated agent skills."

- + Largest catalog (315 plugins, 1,537 skills, 20+ categories)
- + Exceptional educational content (playbooks, tutorials)
- + CCPI CLI for search/install/validate
- Quality varies across 315 plugins
- Discoverability challenge (17,524 files, 132MB)
- No quality ratings or user review system

UI/UX Pro Max Skill

DESIGN INTELLIGENCE SKILL

by nextlevelbuilder • 37.3K ★ • 3,641 forks • 68 UI styles • 96 palettes • v2.2.3 • 18 platforms

"A single-domain design intelligence skill that generates production-ready design systems from natural language. Contains 68 UI styles, 96 color palettes, 57 font pairings, 100+ reasoning rules."

- + 100+ reasoning rules teach Claude why designs work
- + 8-tier priority framework (accessibility → style)
- + Zero Python deps—stdlib only, runs offline
- Single-domain: zero value for non-UI/UX work
- Windows symlink issues (open issue #109)
- CSV data ceiling—no database backend for scale

03

Classification Matrix

All 9 tools compared across the same dimensions for direct, apples-to-apples evaluation.

Table 1: Methodology, Orchestration & Role-Based Tools

DIMENSION	SUPERPOWERS	GSD	GSTACK	COMPOUND ENG.	CLAUDE SKILLS
Category	Methodology	Orchestration	Role-Based	Plugin+CLI	Skills Pack
CC Fit	Native	Native	Native	Native	Native
Setup	Low	Low	Low-Med	Low	Low
Learning Curve	Medium	High	Medium	Med-High	Low
Maintenance	Active	Very Active	Very Active	Very Active	Active
Best For	Process discipline	Complex multi-phase projects	Role-based team workflows	Cross-platform power users	Domain expertise

Table 2: Comprehensive, Design, Marketplace & Autonomous Tools

DIMENSION	EVERYTHING CC	UI/UX PRO MAX	PLUGINS+SKILLS	RALPHY
Category	Comprehensive	Design Intel	Marketplace	Auto Runner
CC Fit	Native	Native	Native	External
Setup	Low-Med	Low	Low-Med	Low
Learning Curve	Medium	Low	Med-High	Low
Maintenance	Active	Active	Active	Moderate
Best For	Standardized team workflows	Design system generation	Plugin catalog + learning	Fire-and-forget automation

Legend

Green = Native / Low / Active Yellow = Medium / Moderate Red = High / External Purple = Category label Blue = Category label

Teal = Category label

04

Deep Comparison

Each criterion is assessed with a named winner, explanation, and trade-offs.

CRITERION	WINNER	RUNNER-UP	TRADE-OFF
Prompt Quality	Superpowers	gstack	Superpowers prevents rationalization; gstack provides role-specific cognitive modes
Agent Architecture	GSD	Compound Eng.	GSD's context-rot solution is novel; CE has more agents
Review Methodology	gstack	Compound Eng.	gstack's Fix-First with suppressions is production-proven; CE has broader swarm
QA Testing	gstack	—	Only tool with persistent headless browser QA. Others rely on MCP or manual testing
Documentation	Claude Skills	Superpowers	CS has deep reference files per skill; SP has exceptional per-skill docs
Domain Coverage	Claude Skills	Everything CC	CS covers 12 domains; ECC covers multiple languages with rules architecture
Ease of Use	Ralphy	UI/UX Pro Max	Ralphy is one command; UI/UX Pro Max is one skill for its domain
Safety	Superpowers	Everything CC	SP enforces TDD gate; ECC has 1,282 security tests with AgentShield

05

Honest Reality Check

No tool is perfect. Here's what each project needs to address honestly.

- **Superpowers** — No parallel execution. Prescriptive workflow can frustrate experienced devs who want to skip phases.
- **GSD** — 32 commands is overwhelming. Rapid release cadence (22 in 2 weeks) risks instability.
- **gstack** — Rails-centric patterns. No built-in orchestration — requires external Conductor for parallelism. Monolithic skill files.
- **Compound Eng.** — 48KB swarm skill strains context. Rails/Ruby bias in review agents. Single maintainer risk.
- **Claude Skills** — Pre-1.0 maturity. No execution engine. Solo contributor.
- **Ralphy** — `--dangerously-skip-permissions` is concerning. Zero quality gates. Bad PRD = bad code fast.
- **Everything CC** — Kitchen sink risk. Hook debugging is opaque. Node.js dependency.
- **Plugins Plus Skills** — Quality varies across 315 plugins. Discoverability is a real challenge.
- **UI/UX Pro Max** — Single domain. CSV backend can't scale.

06

Recommendations by Scenario

SCENARIO	RECOMMENDED STACK	WHY
Solo founder, new project	Superpowers + UI/UX Pro Max	Process discipline + design quality
Team with established codebase	GSD + Claude Skills	Context management + domain expertise
Rapid prototyping	Ralphy	Fire-and-forget speed
Production app with QA needs	gstack	Browser-based QA + role-based review gates
Enterprise / compliance	Everything CC + Superpowers	Security-first + methodology
Learning Claude Code	Plugins Plus Skills	Tutorials + browsable catalog
Cross-platform (multiple AI tools)	Compound Engineering	Only tool that converts between 10 formats

07

Combining Tools

The best setup combines 2–3 tools from different layers, not all 9.

SAFE COMBINATIONS

- Superpowers + Claude Skills** — methodology + domain expertise (different layers, no conflict)
- GSD + Claude Skills** — orchestration + expertise (complementary)
- gstack + Superpowers** — role-based QA/review + methodology (some overlap in review, but different angles)
- Superpowers + UI/UX Pro Max** — methodology + design (orthogonal domains)

RISKY COMBINATIONS

- Superpowers + GSD** — competing workflow philosophies. GSD's commands may conflict with SP's skill triggers.
- Everything CC + Compound Eng.** — both are comprehensive. Combined context footprint is excessive.
- gstack + GSD** — different orchestration models. gstack's monolithic skills vs GSD's lean orchestrator.

Final Ranking

Confidence-scored ranking based on evidence-based evaluation. Scores reflect utility for a production software team using Claude Code as primary development tool.

1	Superpowers	9.2	High	Best prompt engineering. Anti-rationalization is unique. TDD gate is the strongest quality enforcer.
2	Get Shit Done (GSD)	8.8	High	Context rot solution is novel and validated. GSD-2 adds depth. Steep learning curve holds it back.
3	gstack NEW	8.5	Med-High	Role-based cognitive modes are original. Browser QA is unique. Review checklists are production-proven. Rails bias and no orchestration lower universality.
4	Everything Claude Code	8.3	High	Most tested (992 tests). Security-first is rare. Kitchen-sink risk and opaque hooks prevent top ranking.
5	Compound Engineering	8.0	Medium	Cross-platform converter is unique. Deepest feature set. Context footprint and maintainer concentration are concerns.
6	Claude Skills	7.8	Medium	Widest domain coverage. Progressive disclosure is clever. Pre-1.0 and solo contributor limit confidence.
7	UI/UX Pro Max	7.5	Medium	100+ reasoning rules are genuinely educational. Single-domain limits utility.
8	Plugins Plus Skills	7.2	Medium	Best learning resource. Quality variance across 315 plugins is the main risk.
9	Ralphy	6.8	Low	Simplest model but —dangerously—skip—permissions and zero quality gates are hard to recommend.

Note: Rankings reflect utility for a production software team using Claude Code as primary development tool. Your ranking may differ based on your specific needs — a solo frontend developer would rank UI/UX Pro Max much higher.